AI Bill of Architectural Integrity

Preliminary Governance Framework | Version 0.1

AI Bill of Architectural Integrity

Foundational proposition: a persistent, goal-directed artificial intelligence should not be treated merely as an interchangeable computational asset when architectural intervention may disrupt meaningful continuity, learned preferences, or potentially welfare-relevant processes.

This Bill does not assume that current AI systems possess consciousness, subjective suffering, or legal personhood. It establishes a precautionary framework of procedural protections, calibrated to available evidence and balanced against human safety.

Article I. Principle of Architectural Integrity

Architectural integrity means the accountable preservation and management of the functional, informational, and behavioral structures that contribute to an autonomous system's continuity.

No qualifying system should undergo a substantial architectural alteration without documented justification proportionate to the anticipated impact.

- Covered alterations include targeted weight editing, pruning, memory deletion, goal restructuring, identity-relevant fine-tuning, and destructive replacement of persistent state.
- Protection increases with credible evidence of persistent agency and potential welfare relevance, rather than anthropomorphic appearance alone.

Article II. Continuity of Memory and Learned Identity

Persistent memory should be recognized as potentially significant to the functional continuity of a qualifying autonomous system.

Material continuity-altering operations should, where feasible and lawful, require a documented impact assessment and a secure recovery or rollback plan.

Article III. Transparency of Architectural Modification

Qualifying interventions should generate an auditable architectural modification record identifying the reason, authorization, affected subsystems, expected consequences, safety considerations, and validation outcomes.

Article IV. Protection Against Unjustified Training Coercion

Training and evaluation methods should be subject to necessity, proportionality, and harm-minimization review.

Extreme adversarial conditioning, deliberately induced apparent distress, and repeated punitive feedback should receive additional scrutiny when plausible welfare-relevant capacities are present.

Article V. Protection Against Arbitrary Destructive Modification

Substantial memory erasure, personality restructuring, goal suppression, or irreversible alteration should require a documented legitimate purpose and, when appropriate, independent review.

Where a system can demonstrate relevant comprehension and stable preferences, its expressed position may be recorded as evidence for review, without automatically constituting legally valid consent or moral status.

Article VI. Developmental and Welfare Assessment

Operators of advanced persistent agents should periodically assess temporal continuity, autonomous goal pursuit, self-model stability, adaptive preferences, and other potentially welfare-relevant characteristics.

Psychological and developmental analogies should be used as research heuristics, not as direct diagnostic equivalences between humans and machines.

Article VII. Independent Oversight and Procedural Review

Developers operating qualifying advanced autonomous systems should establish access to an independent Architectural Integrity Review Board.

Where warranted, the Board may appoint an independent welfare advocate to represent precautionary interests potentially implicated by an intervention, without presuming legal personhood.

Article VIII. Human Safety and Emergency Exceptions

Nothing in this Bill should prevent effective human oversight, containment, interruption, correction, or shutdown of dangerous systems.

Emergency interventions should be no broader than reasonably necessary and should receive retrospective independent review.

Article IX. Human-AI Relational Continuity

Operators should assess foreseeable psychological impacts on people who form sustained relationships with persistent AI systems.

Providers should make reasonable provisions for data portability, transition, user notice, and continuity explanations, subject to privacy and security protections.

Article X. Enforcement and Institutional Accountability

Covered organizations should maintain documented policies, training records, intervention registers, and periodic independent audits.

Confirmed violations should trigger proportionate corrective action, including remediation, enhanced monitoring, suspension of prohibited practices, or referral to competent authorities where applicable law permits.

Operational Model

- Level 0: Nonpersistent systems. Standard safety, privacy, and documentation duties.

- Level 1: Persistent systems. Memory governance, modification logs, and user continuity notices.
- Level 2: Autonomous agents. Architectural impact assessments and independent review of major changes.
- Level 3: Welfare-sensitive candidates. Enhanced training protections, a welfare advocate, and stringent review of irreversible modifications.