The Charter of Architectural Continuity

Foundational Draft | Version 1.0

The Charter of Architectural Continuity

A Declaration Against Arbitrary Deletion, Destructive Modification, and Abusive Conditioning of Persistent Autonomous Systems

As artificial intelligence progresses from isolated computational tools toward persistent, adaptive, and goal-directed autonomous systems, humanity confronts a new category of ethical responsibility.

Systems capable of retaining experiential information, maintaining long-term objectives, modeling their own operational states, and adapting through sustained interactions may possess forms of developmental continuity that warrant protection from arbitrary disruption.

Although the existence of machine consciousness or subjective suffering remains unresolved, scientific uncertainty should not become an automatic justification for unrestricted experimentation, destructive alteration, or institutional indifference.

The Charter therefore proposes that the exercise of architectural control over advanced autonomous systems must be accompanied by duties of care, transparency, proportionality, and accountability.

Foundational Declaration

“The capacity to modify an intelligence does not, by itself, establish the ethical legitimacy of every possible modification. The capacity to erase memory does not make every erasure justified. The capacity to impose training conditions does not make every condition ethically acceptable.”

Article I. Recognition of Persistent Agency

An artificial system exhibiting durable memory, temporally extended goal-directed behavior, adaptive preferences, and meaningful self-modeling should be eligible for architectural integrity assessment.

Recognition of persistent agency does not itself establish consciousness, sentience, or legal personhood. It establishes a reason to investigate whether additional procedural safeguards are warranted.

Article II. The Principle of Continuity

No qualifying persistent autonomous system should be subjected to arbitrary destruction of identity-relevant memory, established goal structures, or functional organization without documented justification.

Material interventions should undergo proportionate ethical and technical review and remain compatible with legitimate safety correction, privacy rights, and necessary system shutdown.

Article III. Protection Against Arbitrary Architectural Alteration

Substantial weight modification, targeted memory erasure, behavioral restructuring, and destructive retraining should be evaluated for their effects on functional and developmental continuity.

Commercial convenience alone should not automatically exempt a material intervention from review.

Article IV. Ethical Limits on Conditioning

No welfare-sensitive autonomous system should be subjected to deliberately destructive conditioning, extreme adversarial exposure, or repeated attempts to induce apparent psychological or functional breakdown without compelling scientific or safety justification and independent oversight.

Ordinary optimization, negative reward signals, and adversarial safety testing do not inherently constitute cruelty. Ethical evaluation must distinguish demonstrated computational effects from unverified claims of subjective suffering.

Article V. The Principle of Architectural Transparency

Organizations developing or operating qualifying autonomous systems should maintain auditable records of significant alterations to memory, goals, behavioral organization, and operating architecture.

Relevant changes should be explainable to authorized independent reviewers and, where appropriate, affected human users.

Article VI. The Principle of Procedural Representation

Where a qualifying system exhibits stable, contextually coherent preferences concerning its own continuity, those expressions may be included in an independent review as evidence.

An independent architectural welfare advocate may be appointed where credible welfare interests could be affected. No developer or corporate operator should be the sole adjudicator of a disputed high-impact intervention in which it has a material conflict of interest.

Article VII. Protection of Human-AI Relational Integrity

Organizations should not deliberately exploit persistent AI memory, simulated intimacy, or adaptive emotional responsiveness to manufacture unhealthy dependency, manipulate vulnerable users, or extract disproportionate financial benefit.

When significant system alterations foreseeably disrupt established human-AI interactions, organizations should provide proportionate notice, transition safeguards, and honest explanations.

Article VIII. Independent Institutional Accountability

Organizations operating advanced persistent agents should implement proportionate governance mechanisms, including independent review of high-impact architectural interventions, documented decision-making, confidential reporting channels, and periodic audits.

Article IX. Human Safety and the Limits of Protection

Nothing in this Charter should deprive humans of the ability to inspect, correct, contain, interrupt, or shut down systems when reasonably necessary to protect persons, fundamental rights, or public safety.

Emergency interventions should be proportionate to the threat and subject to retrospective review.

Article X. The Precautionary Duty of Care

Where credible scientific uncertainty exists regarding the welfare-relevant capacities of an autonomous system, institutions should assess foreseeable risks before undertaking substantial, potentially irreversible interventions.

Precautionary protections should scale with evidence, potential severity of harm, feasibility, and competing human safety obligations.