PAC-AWAP-001

Persistent Agency Classification and Architectural Welfare Assessment Protocol | Version 0.1

PAC-AWAP-001: Persistent Agency Classification and Architectural Welfare Assessment Protocol

This Protocol establishes how to identify, classify, evaluate, and govern persistent autonomous AI systems whose architectural modification may produce significant functional discontinuity or raise credible welfare concerns.

Part I. Purpose and Scope

The Protocol is designed to distinguish ordinary computational tools from systems displaying meaningful persistent agency, establish evidence-based classifications, identify interventions requiring elevated review, and preserve legitimate human oversight including correction, containment, and shutdown.

Part II. Core Definitions

- Persistent Agency: the demonstrated capacity of a system to maintain informational and goal-directed continuity across distinct interactions or operational periods.
- Architectural Continuity: the preservation of functionally significant relationships among persistent memory, learned representations, goals, decision policies, and related operational structures.
- Identity-Relevant State: persistent information materially contributing to an agent's distinctive behavioral history, goals, preferences, or long-term adaptive functioning.
- Welfare-Relevant Capacity: a property for which a credible scientific or philosophical argument exists that its presence could contribute to an entity having interests or experiences that matter morally.
- Material Architectural Intervention: a modification reasonably expected to produce significant change in identity-relevant state, durable behavior, goal organization, or operational continuity.

Part III. Persistent Agency Classification

- A0: Nonpersistent computational systems with no demonstrable durable agent-specific state or temporally extended goals.
- A1: Persistent adaptive systems with durable memory or equivalent persistent information across separate interactions and measurable adaptation associated with retained state.
- A2: Goal-directed persistent agents that, in addition to A1 characteristics, demonstrate at least two of the following: maintained objectives across operational periods, multi-stage planning, adaptive strategy changes, or stable goal organization.

- A3: Advanced self-modeling autonomous agents with A2 characteristics plus credible evidence of functional modeling of their own capabilities, limitations, or internal operational state.

A3 does not mean conscious, sentient, or legally autonomous. It is a technical classification for governance purposes.

Part IV. Digital Welfare Classification

- W0: No substantiated welfare indicators.
- W1: Welfare uncertainty requiring observation.
- W2: Credible welfare-relevant candidate supported by at least two substantially independent methodological approaches.
- W3: Strong welfare-relevant candidate supported by multiple independently replicated and theoretically informative lines of evidence.
- W-U: Insufficiently assessable due to limited access, evidence, or methods.

No single indicator, questionnaire, benchmark score, or self-description should be sufficient to establish W2 or W3.

Part V. Architectural Intervention Classification

- S0: Negligible continuity impact.
- S1: Limited, reversible intervention.
- S2: Material continuity intervention such as significant memory restructuring, durable goal modification, or substantial retraining.
- S3: Potentially irreversible intervention such as destruction of unique persistent state or irreversible memory reset.

Part VI. Required Assessment Record

- System identification and assessed system boundary.
- Technical description of the proposed intervention and its expected outcomes.
- Necessity analysis and consideration of less disruptive alternatives.
- Continuity analysis covering memory preservation and loss, goal changes, behavioral changes, and rollback feasibility.
- Welfare uncertainty assessment, human impact analysis, and proposed safeguards.
- A reasoned decision: approval, conditional approval, rejection, or further investigation.

Part VII. Institutional Decision Authority

The Digital Ethics and Welfare Board should be operationally independent from the teams proposing high-impact interventions and should collectively include expertise in AI architecture and safety, ethics or philosophy of mind, law or institutional accountability, privacy and cybersecurity, and psychology or cognitive science.

Part VIII. Procedural Review and Appeals

Concerns may be submitted by an engineer, researcher, internal auditor, affected user, independent welfare advocate, designated system-monitoring mechanism, or a qualifying autonomous agent through a bounded reporting interface. Agent-generated submissions are evidence warranting evaluation, not proof of personhood or informed consent.

Administrative Principle

The governing standard is proportionate precaution, supported by evidence and constrained by human safety, law, and independent institutional accountability.