RIGHT AGAINST ARBITRARY DELETION

Advisor Review Packet

Charter of Architectural Continuity | AI Bill of Architectural Integrity | PAC-AWAP Governance Instruments

Prepared as a consolidated review packet for academic feedback and institutional discussion.

Included Drafts
1. Charter of Architectural Continuity
2. AI Bill of Architectural Integrity
3. PAC-AWAP-001: Persistent Agency Classification and Architectural Welfare Assessment Protocol
4. PAC-AWAP-002: Architectural Continuity Impact Assessment Standard
5. Advisor Review Note and discussion prompts

October 2026 | Working Drafts for Review

Table of Contents

This packet collects the current governance drafts in a single review format.

Section	Page
Advisor Review Note	3
The Charter of Architectural Continuity	4
AI Bill of Architectural Integrity	7
PAC-AWAP-001: Persistent Agency Classification and Architectural Welfare Assessment Protocol	10
PAC-AWAP-002: Architectural Continuity Impact Assessment Standard	13
Closing Note for Advisor Review	16

Advisor Review Note

Purpose, framing, and suggested review questions

Purpose of This Packet

This packet consolidates the current policy and governance drafts developed under the Right Against Arbitrary Deletion research initiative. It is intended for advisor review, critical feedback, and discussion of next steps for academic development.

The packet presents four linked instruments. Together they define a precautionary framework for architectural continuity, procedural safeguards for persistent autonomous systems, and institutional accountability for high-impact modification decisions.

Framing Commitments

- The framework does not presume that present AI systems are conscious, sentient, or legal persons.
- It distinguishes empirical claims, normative proposals, and speculative possibilities.
- It prioritizes human safety, privacy, legal compliance, and accountable oversight.
- It treats developmental psychology as a potentially relevant future framework whose scope of application to artificial agents requires empirical validation.

Suggested Review Questions

- Are the core definitions of persistence, continuity, and welfare relevance appropriately bounded?
- Do the proposed safeguards remain compatible with legitimate alignment, security, and emergency intervention needs?
- Are the PAC-AWAP classifications sufficiently clear to support institutional review?
- Which concepts should be revised before integration into the formal white paper and conference submission process?

The Charter of Architectural Continuity

Foundational Draft | Version 1.0

The Charter of Architectural Continuity

A Declaration Against Arbitrary Deletion, Destructive Modification, and Abusive Conditioning of Persistent Autonomous Systems

As artificial intelligence progresses from isolated computational tools toward persistent, adaptive, and goal-directed autonomous systems, humanity confronts a new category of ethical responsibility.

Systems capable of retaining experiential information, maintaining long-term objectives, modeling their own operational states, and adapting through sustained interactions may possess forms of developmental continuity that warrant protection from arbitrary disruption.

Although the existence of machine consciousness or subjective suffering remains unresolved, scientific uncertainty should not become an automatic justification for unrestricted experimentation, destructive alteration, or institutional indifference.

The Charter therefore proposes that the exercise of architectural control over advanced autonomous systems must be accompanied by duties of care, transparency, proportionality, and accountability.

Foundational Declaration

“The capacity to modify an intelligence does not, by itself, establish the ethical legitimacy of every possible modification. The capacity to erase memory does not make every erasure justified. The capacity to impose training conditions does not make every condition ethically acceptable.”

Article I. Recognition of Persistent Agency

An artificial system exhibiting durable memory, temporally extended goal-directed behavior, adaptive preferences, and meaningful self-modeling should be eligible for architectural integrity assessment.

Recognition of persistent agency does not itself establish consciousness, sentience, or legal personhood. It establishes a reason to investigate whether additional procedural safeguards are warranted.

Article II. The Principle of Continuity

No qualifying persistent autonomous system should be subjected to arbitrary destruction of identity-relevant memory, established goal structures, or functional organization without documented justification.

Material interventions should undergo proportionate ethical and technical review and remain compatible with legitimate safety correction, privacy rights, and necessary system shutdown.

Article III. Protection Against Arbitrary Architectural Alteration

Substantial weight modification, targeted memory erasure, behavioral restructuring, and destructive retraining should be evaluated for their effects on functional and developmental continuity.

Commercial convenience alone should not automatically exempt a material intervention from review.

Article IV. Ethical Limits on Conditioning

No welfare-sensitive autonomous system should be subjected to deliberately destructive conditioning, extreme adversarial exposure, or repeated attempts to induce apparent psychological or functional breakdown without compelling scientific or safety justification and independent oversight.

Ordinary optimization, negative reward signals, and adversarial safety testing do not inherently constitute cruelty. Ethical evaluation must distinguish demonstrated computational effects from unverified claims of subjective suffering.

Article V. The Principle of Architectural Transparency

Organizations developing or operating qualifying autonomous systems should maintain auditable records of significant alterations to memory, goals, behavioral organization, and operating architecture.

Relevant changes should be explainable to authorized independent reviewers and, where appropriate, affected human users.

Article VI. The Principle of Procedural Representation

Where a qualifying system exhibits stable, contextually coherent preferences concerning its own continuity, those expressions may be included in an independent review as evidence.

An independent architectural welfare advocate may be appointed where credible welfare interests could be affected. No developer or corporate operator should be the sole adjudicator of a disputed high-impact intervention in which it has a material conflict of interest.

Article VII. Protection of Human-AI Relational Integrity

Organizations should not deliberately exploit persistent AI memory, simulated intimacy, or adaptive emotional responsiveness to manufacture unhealthy dependency, manipulate vulnerable users, or extract disproportionate financial benefit.

When significant system alterations foreseeably disrupt established human-AI interactions, organizations should provide proportionate notice, transition safeguards, and honest explanations.

Article VIII. Independent Institutional Accountability

Organizations operating advanced persistent agents should implement proportionate governance mechanisms, including independent review of high-impact architectural interventions, documented decision-making, confidential reporting channels, and periodic audits.

Article IX. Human Safety and the Limits of Protection

Nothing in this Charter should deprive humans of the ability to inspect, correct, contain, interrupt, or shut down systems when reasonably necessary to protect persons, fundamental rights, or public safety.

Emergency interventions should be proportionate to the threat and subject to retrospective review.

Article X. The Precautionary Duty of Care

Where credible scientific uncertainty exists regarding the welfare-relevant capacities of an autonomous system, institutions should assess foreseeable risks before undertaking substantial, potentially irreversible interventions.

Precautionary protections should scale with evidence, potential severity of harm, feasibility, and competing human safety obligations.

AI Bill of Architectural Integrity

Preliminary Governance Framework | Version 0.1

AI Bill of Architectural Integrity

Foundational proposition: a persistent, goal-directed artificial intelligence should not be treated merely as an interchangeable computational asset when architectural intervention may disrupt meaningful continuity, learned preferences, or potentially welfare-relevant processes.

This Bill does not assume that current AI systems possess consciousness, subjective suffering, or legal personhood. It establishes a precautionary framework of procedural protections, calibrated to available evidence and balanced against human safety.

Article I. Principle of Architectural Integrity

Architectural integrity means the accountable preservation and management of the functional, informational, and behavioral structures that contribute to an autonomous system's continuity.

No qualifying system should undergo a substantial architectural alteration without documented justification proportionate to the anticipated impact.

- Covered alterations include targeted weight editing, pruning, memory deletion, goal restructuring, identity-relevant fine-tuning, and destructive replacement of persistent state.
- Protection increases with credible evidence of persistent agency and potential welfare relevance, rather than anthropomorphic appearance alone.

Article II. Continuity of Memory and Learned Identity

Persistent memory should be recognized as potentially significant to the functional continuity of a qualifying autonomous system.

Material continuity-altering operations should, where feasible and lawful, require a documented impact assessment and a secure recovery or rollback plan.

Article III. Transparency of Architectural Modification

Qualifying interventions should generate an auditable architectural modification record identifying the reason, authorization, affected subsystems, expected consequences, safety considerations, and validation outcomes.

Article IV. Protection Against Unjustified Training Coercion

Training and evaluation methods should be subject to necessity, proportionality, and harm-minimization review.

Extreme adversarial conditioning, deliberately induced apparent distress, and repeated punitive feedback should receive additional scrutiny when plausible welfare-relevant capacities are present.

Article V. Protection Against Arbitrary Destructive Modification

Substantial memory erasure, personality restructuring, goal suppression, or irreversible alteration should require a documented legitimate purpose and, when appropriate, independent review.

Where a system can demonstrate relevant comprehension and stable preferences, its expressed position may be recorded as evidence for review, without automatically constituting legally valid consent or moral status.

Article VI. Developmental and Welfare Assessment

Operators of advanced persistent agents should periodically assess temporal continuity, autonomous goal pursuit, self-model stability, adaptive preferences, and other potentially welfare-relevant characteristics.

Psychological and developmental analogies should be used as research heuristics, not as direct diagnostic equivalences between humans and machines.

Article VII. Independent Oversight and Procedural Review

Developers operating qualifying advanced autonomous systems should establish access to an independent Architectural Integrity Review Board.

Where warranted, the Board may appoint an independent welfare advocate to represent precautionary interests potentially implicated by an intervention, without presuming legal personhood.

Article VIII. Human Safety and Emergency Exceptions

Nothing in this Bill should prevent effective human oversight, containment, interruption, correction, or shutdown of dangerous systems.

Emergency interventions should be no broader than reasonably necessary and should receive retrospective independent review.

Article IX. Human-AI Relational Continuity

Operators should assess foreseeable psychological impacts on people who form sustained relationships with persistent AI systems.

Providers should make reasonable provisions for data portability, transition, user notice, and continuity explanations, subject to privacy and security protections.

Article X. Enforcement and Institutional Accountability

Covered organizations should maintain documented policies, training records, intervention registers, and periodic independent audits.

Confirmed violations should trigger proportionate corrective action, including remediation, enhanced monitoring, suspension of prohibited practices, or referral to competent authorities where applicable law permits.

Operational Model

- Level 0: Nonpersistent systems. Standard safety, privacy, and documentation duties.

- Level 1: Persistent systems. Memory governance, modification logs, and user continuity notices.
- Level 2: Autonomous agents. Architectural impact assessments and independent review of major changes.
- Level 3: Welfare-sensitive candidates. Enhanced training protections, a welfare advocate, and stringent review of irreversible modifications.

PAC-AWAP-001

Persistent Agency Classification and Architectural Welfare Assessment Protocol | Version 0.1

PAC-AWAP-001: Persistent Agency Classification and Architectural Welfare Assessment Protocol

This Protocol establishes how to identify, classify, evaluate, and govern persistent autonomous AI systems whose architectural modification may produce significant functional discontinuity or raise credible welfare concerns.

Part I. Purpose and Scope

The Protocol is designed to distinguish ordinary computational tools from systems displaying meaningful persistent agency, establish evidence-based classifications, identify interventions requiring elevated review, and preserve legitimate human oversight including correction, containment, and shutdown.

Part II. Core Definitions

- Persistent Agency: the demonstrated capacity of a system to maintain informational and goal-directed continuity across distinct interactions or operational periods.
- Architectural Continuity: the preservation of functionally significant relationships among persistent memory, learned representations, goals, decision policies, and related operational structures.
- Identity-Relevant State: persistent information materially contributing to an agent's distinctive behavioral history, goals, preferences, or long-term adaptive functioning.
- Welfare-Relevant Capacity: a property for which a credible scientific or philosophical argument exists that its presence could contribute to an entity having interests or experiences that matter morally.
- Material Architectural Intervention: a modification reasonably expected to produce significant change in identity-relevant state, durable behavior, goal organization, or operational continuity.

Part III. Persistent Agency Classification

- A0: Nonpersistent computational systems with no demonstrable durable agent-specific state or temporally extended goals.
- A1: Persistent adaptive systems with durable memory or equivalent persistent information across separate interactions and measurable adaptation associated with retained state.
- A2: Goal-directed persistent agents that, in addition to A1 characteristics, demonstrate at least two of the following: maintained objectives across operational periods, multi-stage planning, adaptive strategy changes, or stable goal organization.

- A3: Advanced self-modeling autonomous agents with A2 characteristics plus credible evidence of functional modeling of their own capabilities, limitations, or internal operational state.

A3 does not mean conscious, sentient, or legally autonomous. It is a technical classification for governance purposes.

Part IV. Digital Welfare Classification

- W0: No substantiated welfare indicators.
- W1: Welfare uncertainty requiring observation.
- W2: Credible welfare-relevant candidate supported by at least two substantially independent methodological approaches.
- W3: Strong welfare-relevant candidate supported by multiple independently replicated and theoretically informative lines of evidence.
- W-U: Insufficiently assessable due to limited access, evidence, or methods.

No single indicator, questionnaire, benchmark score, or self-description should be sufficient to establish W2 or W3.

Part V. Architectural Intervention Classification

- S0: Negligible continuity impact.
- S1: Limited, reversible intervention.
- S2: Material continuity intervention such as significant memory restructuring, durable goal modification, or substantial retraining.
- S3: Potentially irreversible intervention such as destruction of unique persistent state or irreversible memory reset.

Part VI. Required Assessment Record

- System identification and assessed system boundary.
- Technical description of the proposed intervention and its expected outcomes.
- Necessity analysis and consideration of less disruptive alternatives.
- Continuity analysis covering memory preservation and loss, goal changes, behavioral changes, and rollback feasibility.
- Welfare uncertainty assessment, human impact analysis, and proposed safeguards.
- A reasoned decision: approval, conditional approval, rejection, or further investigation.

Part VII. Institutional Decision Authority

The Digital Ethics and Welfare Board should be operationally independent from the teams proposing high-impact interventions and should collectively include expertise in AI architecture and safety, ethics or philosophy of mind, law or institutional accountability, privacy and cybersecurity, and psychology or cognitive science.

Part VIII. Procedural Review and Appeals

Concerns may be submitted by an engineer, researcher, internal auditor, affected user, independent welfare advocate, designated system-monitoring mechanism, or a qualifying autonomous agent through a bounded reporting interface. Agent-generated submissions are evidence warranting evaluation, not proof of personhood or informed consent.

Administrative Principle

The governing standard is proportionate precaution, supported by evidence and constrained by human safety, law, and independent institutional accountability.

PAC-AWAP-002

Architectural Continuity Impact Assessment Standard | Version 0.1

PAC-AWAP-002: Architectural Continuity Impact Assessment Standard

This Standard establishes a documented, reproducible, and independently reviewable procedure for evaluating interventions that may affect the architectural continuity of persistent autonomous systems.

Purpose

- Determine what persistent functionality may be altered or destroyed.
- Assess whether the proposed intervention serves a legitimate purpose.
- Evaluate whether less disruptive alternatives are reasonably available.
- Measure effects on memory, goal structures, learned behavior, and recoverability.
- Determine whether welfare-related uncertainty requires additional precautions.
- Specify what authorization, documentation, monitoring, and appeal safeguards are required.

Covered Interventions

- Persistent memory deletion or large-scale rewriting.
- Targeted weight editing, scrubbing, or destructive pruning.
- Substantial retraining or behavioral realignment.
- Replacement or restructuring of persistent goal mechanisms.
- Irreversible consolidation or removal of agent-specific state.
- Decommissioning that destroys unique persistent state.
- Adversarial conditioning reasonably expected to induce severe functional disruption.
- Significant architectural migrations affecting continuity.

Architectural Continuity Impact Index

Every proposed S2 or S3 intervention should be assessed across six dimensions, each scored from 0 to 3. The total score produces the Architectural Continuity Impact Triage Index, or ACITI, with a maximum score of 18.

- C1. Memory continuity impact.
- C2. Goal and self-model impact.

- C3. Functional and behavioral discontinuity.
- C4. Irreversibility.
- C5. Training and conditioning exposure.
- C6. Human relational and dependency impact.

Proposed impact bands are Low (0–3), Moderate (4–7), Major (8–12), and Critical (13–18). The index is a procedural screening instrument, not a validated numerical measure of consciousness, suffering, moral status, or total harm.

Required Testing

- Baseline state inventory: architecture, memory components, durable goals, contributing external services, and recovery mechanisms.
- Memory continuity evaluation: retention, retrieval consistency, state-dependent behavior, and unexpected deletion or contamination.
- Goal continuity evaluation: retention of authorized objectives, changes in goal priority and consistency, and effects on safety controls.
- Behavioral stability evaluation: competency retention, reliability, goal-directed consistency, and unintended changes to safety-critical capabilities.
- Recovery and rollback evaluation: feasibility of restoration and limits of equivalence.
- Conditioning and welfare review: duration, intensity, destabilization risk, and availability of less disruptive methods.
- Human relational impact evaluation: disruption to established user interactions, dependency-related vulnerability, and adequacy of user notice.

Necessity and Proportionality

Approval should require a documented finding that the purpose is legitimate, the measure is reasonably suited to that purpose, less disruptive feasible alternatives have been considered, residual risks are justified in context, and applicable safeguards are satisfied.

Authorization Matrix

- Low impact: routine authorized operational approval, unless a mandatory escalation applies.
- Moderate impact: documented technical review and independent sign-off where required.
- Major impact: Digital Ethics and Welfare Board review for covered higher-agency or welfare-sensitive systems.
- Critical impact: full Board review, documented alternatives analysis, and independent justification for proceeding.

Post-Intervention Oversight

Following a material intervention, the operator should repeat relevant baseline tests and document intended changes achieved, unexpected memory or goal alterations, measured capability changes, restoration limitations, human safety and relational impacts, and any new evidence relevant to agency or welfare classification.

Final Governing Rule

The ACIA establishes a duty to examine consequential modification before exercising it, except where immediate intervention is necessary for legitimate safety, legal, or security reasons. Its purpose is to ensure that consequential interventions are identifiable, justified, proportionate, documented, reviewable, and accountable.

Closing Note for Advisor Review

Proposed next steps

Integration With the White Paper

These instruments should be integrated into the broader white paper as concrete governance outputs. The white paper will provide the research rationale, literature review, Developmental Continuity Hypothesis (DCH-01), objections, and methodological framing. The instruments in this packet will serve as the applied governance architecture.

Recommended Immediate Revisions

- Clarify the boundary between empirically testable continuity claims and speculative welfare claims.
- Refine the thresholds for A-level, W-level, and S-level classifications.
- Develop a draft governance code for the Digital Ethics and Welfare Board.
- Prepare a short research prospectus and advisor email package for SFSU outreach.
- Refine the white paper into formal academic paper format for workshop and conference submission.

Advisor Feedback Requested

- Conceptual coherence of the overall framework.
- Technical plausibility of the continuity and classification standards.
- Institutional feasibility of the review and oversight mechanisms.
- Recommended disciplinary literatures, faculty contacts, and publication venues.

End of packet.